

Text Classification Using Transformer Networks (BERT)

Some initialization:

```
In [12]: import random

import torch

import numpy as np

import pandas as pd

from tqdm.notebook import tqdm

# enable tqdm in pandas
tqdm.pandas()

# set to True to use the gpu (if there is one available)
use_gpu = True

# select device
device = torch.device('cuda' if use_gpu and torch.cuda.is_available() else 'cpu')
print(f'device: {device.type}')

# random seed
seed = 1122

# set random seed
if seed is not None:
    print(f'random seed: {seed}')
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
```

device: cuda
random seed: 1122

Read the train/dev/test datasets and create a HuggingFace `Dataset` object:

```
In [13]: def read_data(filename):  
  
    # read csv file  
  
    df = pd.read_csv(filename, header=None)  
  
    # add column names  
  
    df.columns = ['label', 'title', 'description']  
  
    # make labels zero-based  
  
    df['label'] -= 1  
  
    # concatenate title and description, and remove backslashes  
  
    df['text'] = df['title'] + " " + df['description']  
  
    df['text'] = df['text'].str.replace('\\', ' ', regex=False)  
  
    return df  
  
In [14]: labels = open('/kaggle/input/classes-txt/classes.txt').read().splitlines()  
  
    train_df = read_data('/kaggle/input/agnews-pytorch-simple-embed-classif-90/AG_NEWS/train/train.csv')  
    test_df = read_data('/kaggle/input/agnews-pytorch-simple-embed-classif-90/AG_NEWS/test/test.csv')  
    train_df
```

Out[14]:

	label	title	description	text
0	2	Wall St. Bears Claw Back Into the Black (Reuters)	Reuters - Short-sellers, Wall Street's dwindli...	Wall St. Bears Claw Back Into the Black (Reute...
1	2	Carlyle Looks Toward Commercial Aerospace (Reu...	Reuters - Private investment firm Carlyle Grou...	Carlyle Looks Toward Commercial Aerospace (Reu...
2	2	Oil and Economy Cloud Stocks' Outlook (Reuters)	Reuters - Soaring crude prices plus worries\ab...	Oil and Economy Cloud Stocks' Outlook (Reuters...
3	2	Iraq Halts Oil Exports from Main Southern Pipe...	Reuters - Authorities have halted oil export\f...	Iraq Halts Oil Exports from Main Southern Pipe...
4	2	Oil prices soar to all-time record, posing new...	AFP - Tearaway world oil prices, toppling reco...	Oil prices soar to all-time record, posing new...
...
119995	0	Pakistan's Musharraf Says Won't Quit as Army C...	KARACHI (Reuters) - Pakistani President Perve...	Pakistan's Musharraf Says Won't Quit as Army C...
119996	1	Renteria signing a top-shelf deal	Red Sox general manager Theo Epstein acknowle...	Renteria signing a top-shelf deal Red Sox gene...
119997	1	Saban not going to Dolphins yet	The Miami Dolphins will put their courtship of...	Saban not going to Dolphins yet The Miami Dolp...
119998	1	Today's NFL games	PITTSBURGH at NY GIANTS Time: 1:30 p.m. Line: ...	Today's NFL games PITTSBURGH at NY GIANTS Time...
119999	1	Nets get Carter from Raptors	INDIANAPOLIS -- All-Star Vince Carter was trad...	Nets get Carter from Raptors INDIANAPOLIS -- A...

120000 rows × 4 columns

In [15]:

```
from sklearn.model_selection import train_test_split

train_df, eval_df = train_test_split(train_df, train_size=0.9)

train_df.reset_index(inplace=True, drop=True)
eval_df.reset_index(inplace=True, drop=True)

print(f'train rows: {len(train_df.index):,}')
print(f'eval rows: {len(eval_df.index):,}')
print(f'test rows: {len(test_df.index):,}')

train rows: 108,000
eval rows: 12,000
test rows: 7,600
```

In [16]:

```
from datasets import Dataset, DatasetDict
```

```
ds = DatasetDict()

ds['train'] = Dataset.from_pandas(train_df)

ds['validation'] = Dataset.from_pandas(eval_df)

ds['test'] = Dataset.from_pandas(test_df)

ds
```

```
Out[16]: DatasetDict({
  train: Dataset({
    features: ['label', 'title', 'description', 'text'],
    num_rows: 108000
  })
  validation: Dataset({
    features: ['label', 'title', 'description', 'text'],
    num_rows: 12000
  })
  test: Dataset({
    features: ['label', 'title', 'description', 'text'],
    num_rows: 7600
  })
})
```

Tokenize the texts:

```
In [17]: from transformers import AutoTokenizer
```

```
transformer_name = 'bert-base-cased'

tokenizer = AutoTokenizer.from_pretrained(transformer_name)
```

```
/opt/conda/lib/python3.10/site-packages/transformers/tokenization_utils_base.py:1617:
FutureWarning: `clean_up_tokenization_spaces` was not set. It will be set to `True` b
y default. This behavior will be deprecated in transformers v4.45, and will be then s
et to `False` by default. For more details check this issue: https://github.com/huggi
ngface/transformers/issues/31884
warnings.warn(
```

Esta parte del código usa la librería Transformers para cargar el tokenizer del modelo preentrenado de BERT: bert-base-cased. Lo que hace es que el tokenizer convierte el texto en pedacitos (tokens) que el modelo puede entender. Con AutoTokenizer.from_pretrained, se descarga el tokenizer y así se puede preparar el texto para que el modelo lo procese.

```
In [18]: def tokenize(examples):

    return tokenizer(examples['text'], truncation=True)

train_ds = ds['train'].map(

    tokenize, batched=True,
```

```
Map: 0%|          | 0/108000 [00:00<?, ? examples/s]
Map: 0%|          | 0/12000 [00:00<?, ? examples/s]
```

	label	input_ids	token_type_ids	attention_mask
0	2	[101, 16752, 13335, 1186, 2101, 6690, 9717, 11...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
1	1	[101, 145, 11680, 17308, 9741, 2428, 150, 1469...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
2	2	[101, 1418, 14099, 27086, 1494, 1114, 4031, 11...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
3	1	[101, 2404, 117, 6734, 1996, 118, 1565, 5465, ...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
4	3	[101, 142, 10044, 27302, 4317, 1584, 3273, 111...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
...
107995	1	[101, 4922, 2274, 1654, 1112, 10503, 1505, 112...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
107996	3	[101, 10605, 24632, 11252, 21285, 10221, 118, ...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
107997	2	[101, 13832, 3484, 11300, 4060, 5058, 112, 188...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
107998	3	[101, 142, 13675, 3756, 5795, 2445, 1104, 109,...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]
107999	2	[101, 157, 16450, 1658, 5302, 185, 7776, 11006...]	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...]	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...]

Este código aplica el tokenizer al dataset y lo prepara para entrenar el modelo. La función `tokenize` toma como entrada ejemplos del dataset y les aplica el tokenizer para evitar que los textos sean demasiados largos. Primero, se aplica la función `tokenize` a todos los ejemplos en grupo del dataset de entrenamiento. Además, elimina las columnas 'title', 'description' y 'text' del dataset, dejando solo los datos necesarios después de tokenizar. Luego, se hace lo mismo para el conjunto de validación.

Create the transformer model:

```
In [19]: from torch import nn

from transformers.modeling_outputs import SequenceClassifierOutput

from transformers.models.bert.modeling_bert import BertModel, BertPreTrainedModel

# https://github.com/huggingface/transformers/blob/65659a29cf5a079842e61a63d57fa244742

class BertForSequenceClassification(BertPreTrainedModel):

    def __init__(self, config):

        super().__init__(config)

        self.num_labels = config.num_labels

        self.bert = BertModel(config)

        self.dropout = nn.Dropout(config.hidden_dropout_prob)

        self.classifier = nn.Linear(config.hidden_size, config.num_labels)

        self.init_weights()

    def forward(self, input_ids=None, attention_mask=None, token_type_ids=None, labels=None):

        outputs = self.bert(

            input_ids,

            attention_mask=attention_mask,

            token_type_ids=token_type_ids,

            **kwargs,

        )

        cls_outputs = outputs.last_hidden_state[:, 0, :]

        cls_outputs = self.dropout(cls_outputs)

        logits = self.classifier(cls_outputs)

        loss = None

        if labels is not None:

            loss_fn = nn.CrossEntropyLoss()
```

```

        loss = loss_fn(logits, labels)

    return SequenceClassifierOutput(

        loss=loss,

        logits=logits,

        hidden_states=outputs.hidden_states,

        attentions=outputs.attentions,

    )

```

Este código define una clase llamada **BertForSequenceClassification** que adapta el modelo **BERT** para clasificar texto. Toma un texto como entrada, lo procesa con **BERT** para extraer las representaciones del texto y luego utiliza una capa lineal final para clasificar el texto en cierta categoría.

Primero, en la inicialización, se configura el modelo con:

- Una instancia de **BertModel**, que es la parte que entiende el texto.
- Una capa de *dropout* para evitar que el modelo aprenda patrones que no generalizan bien.
- Una capa lineal que toma las representaciones de BERT y las convierte en las probabilidades de cada clase.

En el método **forward**, el texto pasa por BERT, se toma el primer token (**[CLS]**) como resumen de toda la secuencia, y luego se aplica el *dropout* y la capa lineal para generar las predicciones. Si se incluyen etiquetas (las respuestas correctas), también calcula la amiento.

Finalmente, devuelve un objeto que incluye las predicciones, li aplica), los estados ocultos y las atecíficas.

```

In [20]: from transformers import AutoConfig

config = AutoConfig.from_pretrained(

    transformer_name,

    num_labels=len(labels),

)

model = (

    BertForSequenceClassification

    .from_pretrained(transformer_name, config=config)

)

```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-cased and are newly initialized: ['classifier.bias', 'classifier.weight']
 You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

En este código, con `AutoConfig`, se obtiene la configuración del modelo preentrenado especificado por `transformer_name` y se ajusta el número de etiquetas para que coincida con la cantidad de categorías en el problema. Luego, con `BertForSequenceClassification.from_pretrained`, se carga el modelo preentrenado de **Hugging Face**, utilizando la configuración que se definios.

Create the trainer object and train:

```
In [21]: from transformers import TrainingArguments


num_epochs = 2

batch_size = 24

weight_decay = 0.01

model_name = f'{transformer_name}-sequence-classification'

training_args = TrainingArguments(
    output_dir=model_name,
    log_level='error',
    num_train_epochs=num_epochs,
    per_device_train_batch_size=batch_size,
    per_device_eval_batch_size=batch_size,
    evaluation_strategy='epoch',
    weight_decay=weight_decay,
)
```

/opt/conda/lib/python3.10/site-packages/transformers/training_args.py:1545: FutureWarning: `evaluation_strategy` is deprecated and will be removed in version 4.46 of  Transformers. Use `eval_strategy` instead
 warnings.warn(

En este código se define el número de épocas, el tamaño del batch y la regularización para evitar el sobreajuste.

Con `TrainingArguments`, se establecen configuraciones clave como:

- **output_dir** : La carpeta donde se guardarán los modelos entrenados y los resultados.
- **num_train_epochs** : El número de veces que el modelo verá el conjunto completo de entrenamiento (2 épocas en este caso).
- **per_device_train_batch_size** y **per_device_eval_batch_size** : El tamaño de los lotes para entrenamiento y evaluación.
- **evaluation_strategy** : Se configura para evaluar el modelo al final de cada época.
- **log_level** : Solo se mostrarán errores durante el entrenamiento.

```
In [22]: from sklearn.metrics import accuracy_score

def compute_metrics(eval_pred):
    y_true = eval_pred.label_ids
    y_pred = np.argmax(eval_pred.predictions, axis=-1)
    return {'accuracy': accuracy_score(y_true, y_pred)}
```

Este código define una función llamada **compute_metrics**, que se encarga de calcular el accuracy de las predicciones realizadas. Toma como entrada **eval_pred**, un objeto que contiene:

- **label_ids** : Las etiquetas reales del conjunto de evaluación.
- **predictions** : Las predicciones del modelo.

Dentro de la función, se usa **np.argmax** para convertir las probabilidades en predicciones finales (seleccionando la clase con mayor probabilidad). Luego, compara estas predicciones con las etiquetas reales usando **accuracy_score** de **scikit-learn**, que calcula el porcentaje de aciertos.

acción.

In [23]: `from transformers import Trainer`

```
trainer = Trainer(
    model=model,
    args=training_args,
    compute_metrics=compute_metrics,
    train_dataset=train_ds,
    eval_dataset=eval_ds,
    tokenizer=tokenizer,
)
```

Este código crea un objeto Trainer, para entrenar y evaluar modelos de manera fácil y eficiente. Aquí se integran todos los elementos necesarios para que el entrenamiento funcione correctamente:

- `model`: El modelo que se va a entrenar, en este caso, BERT adaptado para clasificación de secuencias.
- `args`: Los argumentos del entrenamiento definidos previamente en `training_args`, como el número de épocas, el tamaño de los lotes, y cómo manejar la evaluación.
- `compute_metrics`: Una función que calcula métricas durante la evaluación. Aquí se utiliza `compute_metrics` para calcular la precisión (accuracy).
- `train_dataset` y `eval_dataset`: Los datos de entrenamiento (`train_ds`) y evaluación (`eval_ds`) que el modelo utilizará para aprender y validarse.
- `tokenizer`: El tokenizer que transforma el texto en tokens, asegurando que los datos se procesen correctamente antes de entrar al modelo.

Con este Trainer, puedes entrenar el modelo, evaluarlo y ajustar sus parámetros automáticamente, simplificando mucho el proceso y evitando errores manuales.

In [24]: `trainer.train()`

```
wandb: WARNING The `run_name` is currently set to the same value as `TrainingArgument
s.output_dir`. If this was not intended, please specify a different run name by setti
ng the `TrainingArguments.run_name` parameter.
wandb: Using wandb-core as the SDK backend. Please refer to https://wandb.me/wandb-co
re for more information.
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wand
b.me/wandb-server)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize
wandb: Paste an API key from your profile and hit enter, or press ctrl+c to quit:
```

```
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
VBox(children=(Label(value='Waiting for wandb.init()...\r'), FloatProgress(value=0.01
1112841188888625, max=1.0...
```

Tracking run with wandb version 0.18.3
Run data is saved locally in /kaggle/working/wandb/run-20241124_071032-60mbfb0g
Syncing run **bert-base-cased-sequence-classification** to [Weights & Biases \(docs\)](#)
View project at <https://wandb.ai/mansoor35/huggingface>
View run at <https://wandb.ai/mansoor35/huggingface/runs/60mbfb0g>

```
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.  
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):  
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.  
    warnings.warn('Was asked to gather along dimension 0, but all ')
```

[4500/4500 55:38, Epoch 2/2]

Epoch	Training Loss	Validation Loss	Accuracy
1	0.189400	0.170879	0.941833
2	0.102300	0.163020	0.946250

```

/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.
    warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.
    warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.
    warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.
    warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.
    warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: FutureWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.autocast('cuda', args...)` instead.
    with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocast_enabled):
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarning: Was asked to gather along dimension 0, but all input tensors were scalars; will instead unsqueeze and return a vector.
    warnings.warn('Was asked to gather along dimension 0, but all '

```

```
ng: Was asked to gather along dimension 0, but all input tensors were scalars; will i
nstead unsqueeze and return a vector.
```

```
warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: Futur
eWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.aut
ocast('cuda', args...)` instead.
```

```
with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocas
t_enabled):
```

```
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarni
ng: Was asked to gather along dimension 0, but all input tensors were scalars; will i
nstead unsqueeze and return a vector.
```

```
warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: Futur
eWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.aut
ocast('cuda', args...)` instead.
```

```
with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocas
t_enabled):
```

```
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarni
ng: Was asked to gather along dimension 0, but all input tensors were scalars; will i
nstead unsqueeze and return a vector.
```

```
warnings.warn('Was asked to gather along dimension 0, but all '
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/parallel_apply.py:79: Futur
eWarning: `torch.cuda.amp.autocast(args...)` is deprecated. Please use `torch.amp.aut
ocast('cuda', args...)` instead.
```

```
with torch.cuda.device(device), torch.cuda.stream(stream), autocast(enabled=autocas
t_enabled):
```

```
/opt/conda/lib/python3.10/site-packages/torch/nn/parallel/_functions.py:68: UserWarni
ng: Was asked to gather along dimension 0, but all input tensors were scalars; will i
nstead unsqueeze and return a vector.
```

```
warnings.warn('Was asked to gather along dimension 0, but all ')
```

```
Out[24]: TrainOutput(global_step=4500, training_loss=0.16224037170410155, metrics={'train_run
ime': 3348.6705, 'train_samples_per_second': 64.503, 'train_steps_per_second': 1.344,
'total_flos': 1.5600315493990656e+16, 'train_loss': 0.16224037170410155, 'epoch': 2.
0})
```

En esta parte del código se entrena el modelo.

Evaluate on the test partition:

```
In [25]: test_ds = ds['test'].map(

    tokenize,

    batched=True,

    remove_columns=['title', 'description', 'text'],

)

test_ds.to_pandas()
```

```
Map: 0%|          | 0/7600 [00:00<?, ? examples/s]
```

	label	input_ids	token_type_ids	attention_mask
0	2	[101, 11284, 1116, 1111, 157, 151, 12966, 1170...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
1	3	[101, 1109, 6398, 1110, 1212, 131, 2307, 7219,...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
2	3	[101, 148, 1183, 119, 1881, 16387, 1116, 4468,...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
3	3	[101, 11689, 15906, 6115, 12056, 1116, 1370, 2...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
4	3	[101, 11917, 8914, 119, 19294, 4206, 1106, 215...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
...
7595	0	[101, 5596, 1103, 1362, 5284, 5200, 3234, 1384...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
7596	1	[101, 159, 7874, 1110, 2709, 1114, 13875, 1556...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
7597	1	[101, 16247, 2972, 9178, 2409, 4271, 140, 1418...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
7598	2	[101, 126, 1104, 1893, 8167, 10721, 4420, 1107...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
7599	2	[101, 142, 2064, 4164, 3370, 1154, 13519, 1116...	[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...

Se aplica el mismo proceso al conjunto de prueba.

output

```
Out[26]: PredictionOutput(predictions=array([[ 0.36316976, -4.0587187 ,  4.708022 , -1.248963
 2 ],
      [ 0.23322392, -3.313729 , -3.7166765 ,  5.9260798 ],
      [ 0.5275277 , -3.135042 , -3.806155 ,  5.569708 ],
      ...,
      [-1.4069046 ,  7.2616873 , -2.2133346 , -3.5286772 ],
      [-0.6832599 , -3.6432836 ,  6.008366 , -2.3655388 ],
      [-3.0042052 , -3.8196225 ,  3.4429522 ,  2.3862102 ]],
      dtype=float32), label_ids=array([2, 3, 3, ..., 1, 2, 2]), metrics={'test_loss':
0.17101630568504333, 'test_accuracy': 0.95, 'test_runtime': 37.9558, 'test_samples_per_second': 200.233, 'test_steps_per_second': 4.189})
```

Se predicen los resultados.

```
In [27]: from sklearn.metrics import classification_report
```

```
y_true = output.label_ids
```

```
y_pred = np.argmax(output.predictions, axis=-1)
```

```
target_names = labels
```

```
print(classification_report(y_true, y_pred, target_names=target_names))
```

	precision	recall	f1-score	support
World	0.97	0.96	0.96	1900
Sports	0.99	0.99	0.99	1900
Business	0.93	0.91	0.92	1900
Sci/Tech	0.91	0.94	0.93	1900
accuracy			0.95	7600
macro avg	0.95	0.95	0.95	7600
weighted avg	0.95	0.95	0.95	7600

Se muestra el desempeño del modelo.